

### Iteration 3 (DL methods)

#### **Problem Statement:**

To find the actions of the JCB arm, i.e; (idle, dig, swing, dump)

#### **Method 1**

##### **CNN-with-regression** (doesn't work as expected)

- **Model:** 5-layerd 2d-cnn with a mean square error loss over logits. Each layer has a 2d convolution layer with relu activation followed by 2d max-pooling.
- **Output:** vector of length 6, each pair represent the location of one of the three joint points. (3\*2)
- **Trained:** >=1000 epochs
- **Dataset:** Train data: 40 frames, manually labelled all three points on each image  
Test data: 20 frames (unable to create large dataset)

#### **Cons:**

- x Loss was not able to converge to an acceptable range.
- ✓ Faster training time

**Note:** The plan was to use these joint points and train a time-series based RNN network. Since due to lack of enough training data, the loss was not able to converge.

#### **Method 2.1**

##### **3D-CNN v1** (works, but slow)

- **Model:** 6-layerd 3d-cnn with a cross entropy loss over softmax-logits. Each layer has a 3d convolution layer (3 frames – 3 filters) with relu activation followed by 3d max-pooling.
- **Output:** one-hot vector of actions
- **Trained:** 12 epochs with k-fold training
- **Dataset:** Train data: used first 1:40 sec for training (approx 200 frames)  
Test data: tested the model on remaining part of the video

#### **Cons:**

- x High training time
- x More no.of epochs needed
- ✓ Higher accuracy

#### **Method 2.2**

##### **3D-CNN v2** (Final working model !)

- **Model:** 6-layerd 3d-cnn with a cross entropy loss over softmax-logits. Each layer has a 3d convolution layer (2 frames – 1 filters) with relu activation followed by 3d max-pooling.
- **Output:** one-hot vector of actions

- **Trained:** 12 epochs with k-fold training
- **Dataset:** Train data: used first 1:40 sec for training (approx 200 frames)  
Test data: tested the model on remaining part of the video

**Pros:** (when compared to previous model - v1)

- ✓ Relative drop in training time
- ✓ Less no.of training epochs needed
- ✓ Faster processing rate during testing
- ✓ Near to higher accuracy

**Results:**

**output video link :** <https://www.youtube.com/watch?v=ZFAAfcEaXpM>

(Please play at fullscreen HD resolution, or the labels might not be visible.)

(Actions being detected are – idle, digging, swing, dumping)

**Note:**

Also tested on a similar video from youtube, **without any additional training**, and it was able to detect digging and swings, for most of the time. (If trained, will perform better.)

**output video link :** <https://www.youtube.com/watch?v=cPDQ7AMzt2k>

**Processing time:**

Video frame rate is 24 frames per sec. The model process 12 frames at a time.

On my machine (i5, 8GB RAM, NO GPU) it takes approx 1.5 sec to process 12 frames.

On google colab (12GB GPU Memory) it takes approx 0.01 sec to process 12 frames.

**Screenshots of runtime in Google colab**

Training phase in colab:

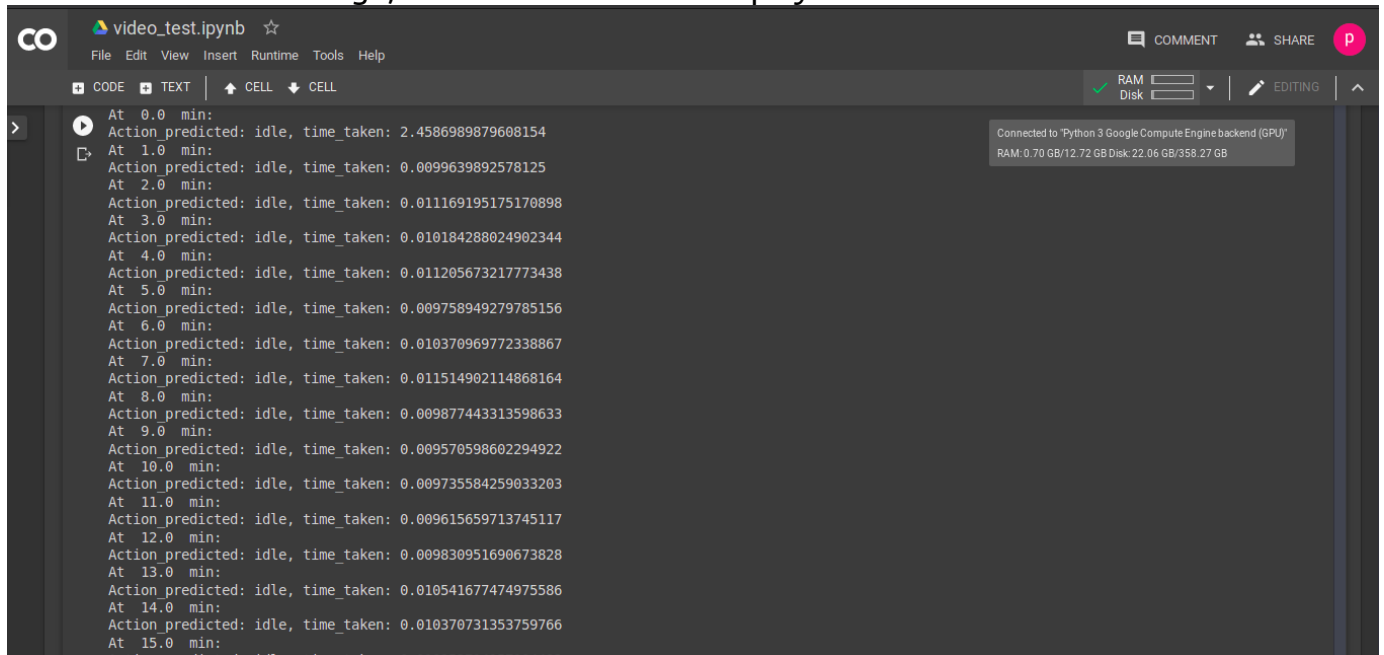
```

cnn_to_actions_multi_img_v3.ipynb
File Edit View Insert Runtime Tools Help
CODE TEXT CELL CELL
CONNECT EDITING
[ ] WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/training/saver.py:1266: checkpoint_exists (from tensorflow.python.trai
Instructions for updating:
Use standard file APIs to check for files with this prefix.
INFO:tensorflow:Restoring parameters from checkpoints/convnet_layers/cnn_to_act_single_img-1995
Loss at step 1999: 0.006235549226403236
Loss at step 2019: 0.014193419367074966
Loss at step 2039: 0.1944298893213272
Loss at step 2059: 0.002012729411944747
Loss at step 2079: 3.814689989667386e-06
Average loss at epoch 0: 0.10438452373877226
Took: 187.81116104125977 seconds
Accuracy at epoch 0: 0.92
Took: 16.09071660041809 seconds
Loss at step 2099: 0.022937802597880363
Loss at step 2119: 0.04116762429475784
Loss at step 2139: 0.2185090333223343
Loss at step 2159: 0.034027110785245895
Loss at step 2179: 0.0
Average loss at epoch 1: 0.05879114313778551
Took: 175.03471207618713 seconds
Accuracy at epoch 1: 1.0
Took: 13.707718133926392 seconds
Loss at step 2199: 0.00015710550360381603
Loss at step 2219: 0.00026067672297358513
Loss at step 2239: 0.195296049118042
Loss at step 2259: 0.01831664890050888
Loss at step 2279: 0.0
Average loss at epoch 2: 0.02481815965536228
Took: 182.9464852809906 seconds
Accuracy at epoch 2: 1.0
Took: 14.492927074432373 seconds

```

Testing phase in colab: (takes around 0.01 sec to process 12 frames)

Note: Also in the image, the GPU allocation is displayed



The screenshot shows a Jupyter Notebook titled 'video\_test.ipynb' with a menu bar (File, Edit, View, Insert, Runtime, Tools, Help) and a toolbar with icons for CODE, TEXT, CELL, and a play button. The notebook contains a series of code cells, each followed by an output cell. The output cells display the following text:

```
At 0.0 min:  
Action_predicted: idle, time_taken: 2.4586989879608154  
At 1.0 min:  
Action_predicted: idle, time_taken: 0.0099639892578125  
At 2.0 min:  
Action_predicted: idle, time_taken: 0.011169195175170898  
At 3.0 min:  
Action_predicted: idle, time_taken: 0.010184288024902344  
At 4.0 min:  
Action_predicted: idle, time_taken: 0.011205673217773438  
At 5.0 min:  
Action_predicted: idle, time_taken: 0.009758949279785156  
At 6.0 min:  
Action_predicted: idle, time_taken: 0.010370969772338867  
At 7.0 min:  
Action_predicted: idle, time_taken: 0.011514902114868164  
At 8.0 min:  
Action_predicted: idle, time_taken: 0.009877443313598633  
At 9.0 min:  
Action_predicted: idle, time_taken: 0.009570598602294922  
At 10.0 min:  
Action_predicted: idle, time_taken: 0.009735584259033203  
At 11.0 min:  
Action_predicted: idle, time_taken: 0.009615659713745117  
At 12.0 min:  
Action_predicted: idle, time_taken: 0.009830951690673828  
At 13.0 min:  
Action_predicted: idle, time_taken: 0.010541677474975586  
At 14.0 min:  
Action_predicted: idle, time_taken: 0.010370731353759766  
At 15.0 min:  
Action_predicted: idle, time_taken: 0.010370731353759766
```

On the right side of the notebook, there is a status bar showing 'Connected to "Python 3 Google Compute Engine backend (GPU)"' and 'RAM: 0.70 GB/12.72 GB Disk: 22.06 GB/358.27 GB'. There are also buttons for 'COMMENT', 'SHARE', and 'EDITING'.

### **Further improvements:**

- To create a generalized model, so that given any new video with different point-of-view / environment setting our algorithm must be able to detect the actions, with minimal amount of training.
- Conversion of fish-eye camera view to a flat image using intrinsic and extrinsic camera parameters for viewing angles.

### **Few ideas for scope of improvement:**

- unet(segmentation) + cnn encoder (feature extraction) + rnn (time series prediction)